

# OPERATING SYSTEM

## Process Management

# Chapter 2

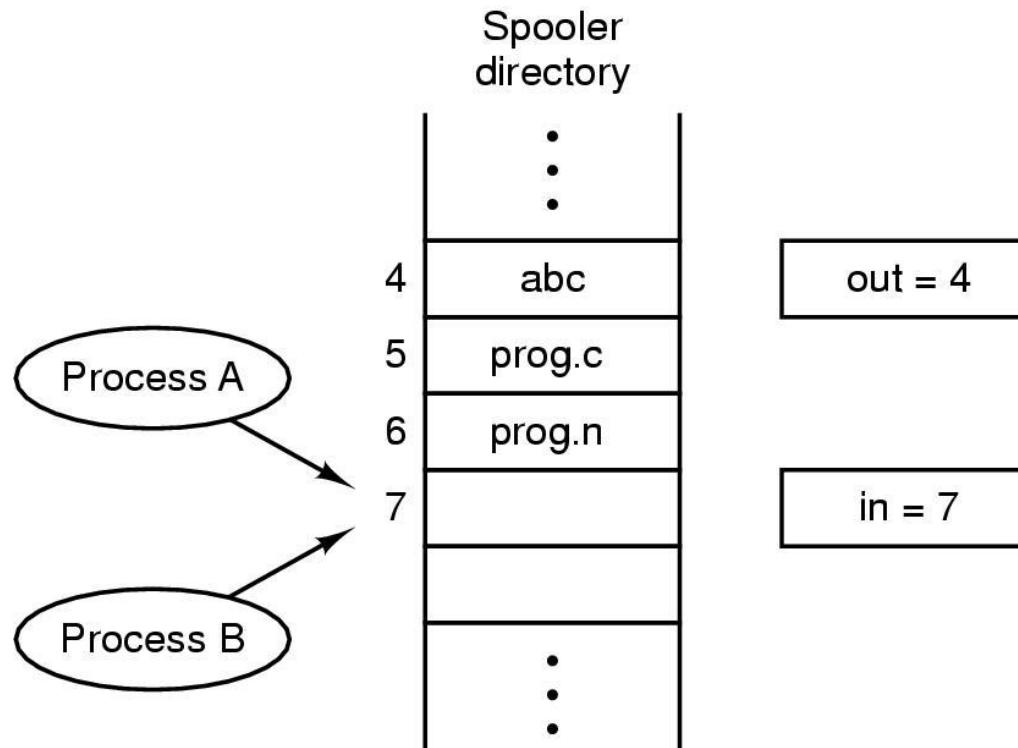
## Processes and Threads

### 2.4 Interprocess communication

# Interprocess Communication

- Processes often need to communicate with other processes
- Example: shell pipeline
- Now we'll look at some issues related to IPC inshaAllah.

# Race Conditions



Two processes want to access shared memory at same time

# How to avoid race condition?

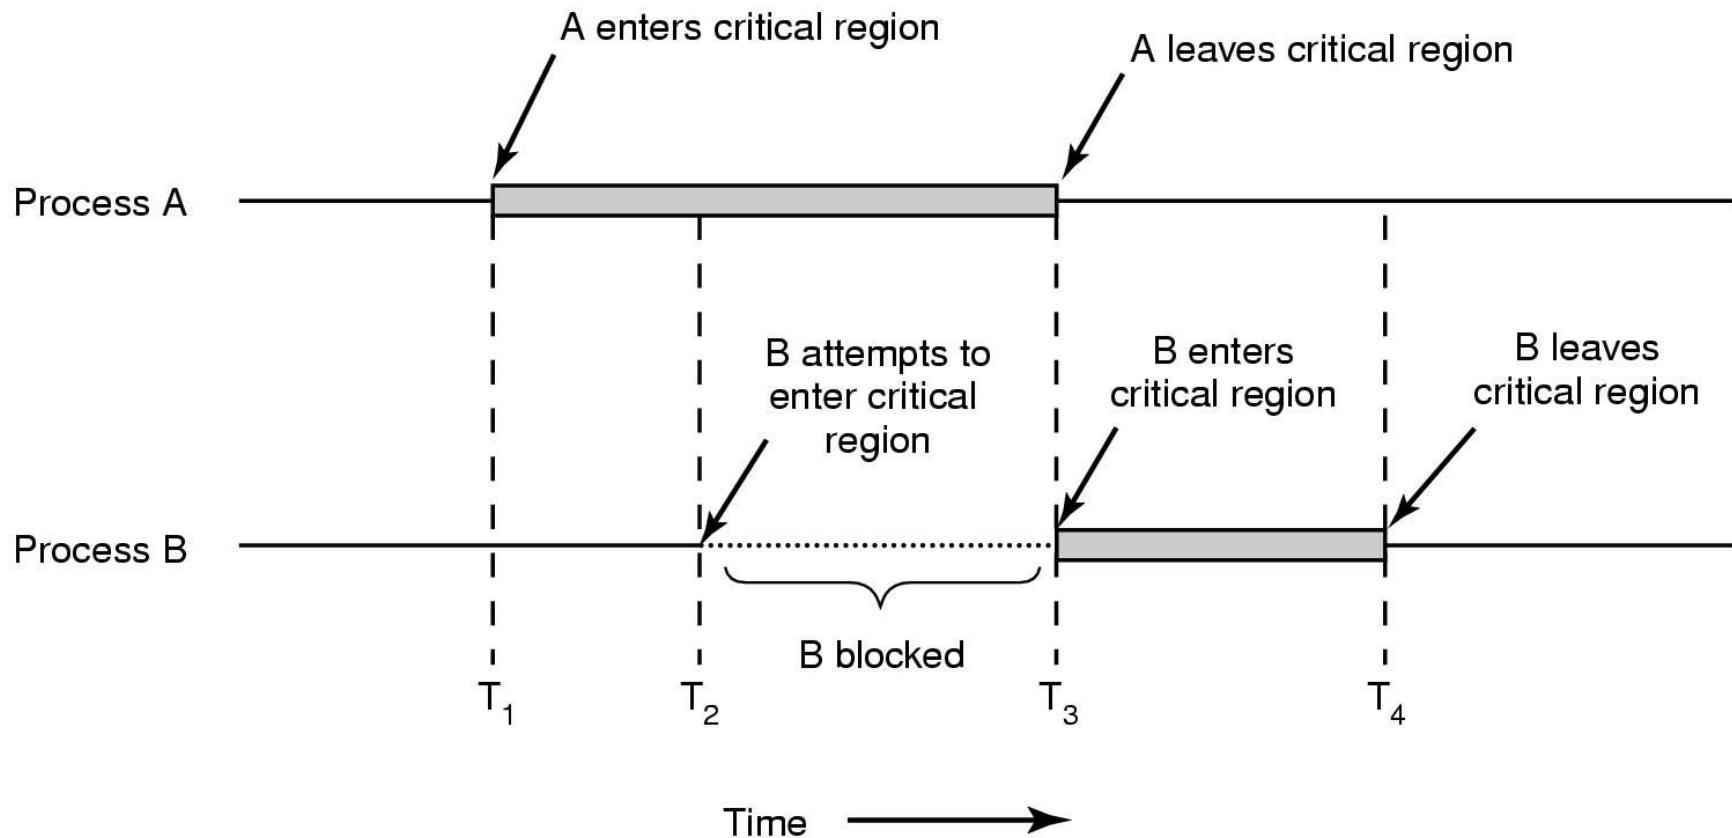
- Mutual Exclusion:
  - No 2 processes use a shared data at the same time
  - If one process is using a shared data , the other process will be excluded from doing the same thing
- Critical Region:
  - That part of the program that do critical things such as accessing shared memory
  - Can lead to race condition

# Solution

## Four conditions for a good solution

1. No two processes simultaneously in critical region
2. No assumptions made about speeds or numbers of CPUs
3. No process running outside its critical region may block another process
4. No process must wait forever to enter its critical region

# Solution



Mutual exclusion using critical regions

# Mutual Exclusion with Busy Waiting (solution 1)

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region();  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region();  
}
```

(b)

Proposed solution to critical region problem

(a) Process 0.

(b) Process 1.

# Problem

- Busy waiting: Continuously testing a variable until some value appear
  - Wastes CPU time
  - Strict Alternation of two process
- Violates condition 3
  - When one process is much slower than the other

# Problem

## Scenario of Violation 3 (Step-by-Step)

**Violation:** P0 is blocked by P1, which is not in its critical region.

**Reason:** P1 is much slower than P0

Step	Process 0	Process 1	Value of turn	Situation
1	P0 enters CR	P1 waiting (turn=0)	0	OK
2	P0 leaves CR, sets turn=1	-	1	P1 can enter now
3	P0 Waiting	P1 quickly enters CR, then exits and sets turn=0	0	OK
4	P0 leaves CR, sets turn=1	P1 still in non-CR	turn = 1	
5	P0 finish non-CR fast → wants to enter CR again	P1 still in non-CR	turn = 1	P0 blocked in while(turn != 0) even though P1 is not in CR

# Mutual Exclusion with Busy Waiting (solution 2)

- Dekker's Algorithm: Dutch mathematician T. Dekker first developed a *software-based solution* to the mutual exclusion problem that did not require strict alternation, combining ideas of *turn-taking* and *lock/warning variables*.
- Peterson's Simplification: In 1981, G. L. Peterson introduced a simpler and more elegant algorithm for achieving mutual exclusion, effectively replacing Dekker's earlier solution.
- Mechanism Overview: In Peterson's algorithm, each process calls `enter_region()` before entering its critical section and `leave_region()` after finishing, ensuring safe and cooperative access to shared resources.

# Mutual Exclusion with Busy Waiting (solution 2)

```
#define FALSE 0
#define TRUE 1
#define N      2          /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process); /* process is 0 or 1 */
{
    int other;               /* number of the other process */

    other = 1 - process;     /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;          /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Peterson's solution for achieving mutual exclusion

# Mutual Exclusion with Busy Waiting

## Peterson's Solution: Analysis(1)

Step	Action / Situation	Explanation / Result
1	Both processes are outside their critical regions.	<code>interested[0] = FALSE, interested[1] = FALSE.</code> No one is using shared data yet.
2	Process 0 calls <code>enter_region(0)</code> .	Sets <code>interested[0] = TRUE</code> and <code>turn = 0</code> . Since P1 is not interested, P0 enters its critical region immediately.
3	Process 1 tries to enter later.	P1 sees <code>interested[0] = TRUE</code> , so it waits until P0 finishes.
4	Both processes call <code>enter_region()</code> almost at the same time.	Both set <code>interested[ ] = TRUE</code> and each sets <code>turn</code> to its own number. The process that writes to <code>turn</code> last decides the value.
5	Suppose <code>turn = 1</code> (P1 wrote last).	P0's condition ( <code>turn == 0</code> ) is false, so P0 enters critical region. P1 keeps waiting because it's P0's turn.
6	P0 exits critical region.	P0 sets <code>interested[0] = FALSE</code> . Now P1's waiting condition becomes false.
7	P1 enters its critical region.	Safe alteration occurs; mutual exclusion is preserved.

# Mutual Exclusion with Busy Waiting

## Peterson's Solution: Analysis(2)

- Let Process 1 is not interested and Process 0 calls `enter_region` with 0
- So,  $\text{turn} = 0$  and  $\text{interested}[0] = \text{true}$  and Process 0 is in CR
- Now if Process 1 calls `enter_region`, it will hang there until  $\text{interested}[0]$  is false. Which only happens when Process 0 calls `leave_region` i.e. leaves the CR

# Mutual Exclusion with Busy Waiting

## Peterson's Solution: Analysis(3)

- Let both processes call `enter_region` simultaneously
- Say `turn = 1`. (i.e. Process 1 stores last)
- Process 0 enters critical region: `while (turn == 0 && ...)` returns false since `turn = 1`.
- Process 1 loops until process 0 exits: `while (turn == 1 && interested[0] == true)` returns true.
- Done!!

# Mutual Exclusion with Busy Waiting: TSL

- Requires hardware support
- TSL instruction: test and set lock
  - Reads content of *lock* into a Register
  - Stores a nonzero value at *lock*.
- CPU executing TSL locks the memory prohibiting other CPUs



Indivisible/Atomic

enter\_region:

```
TSL REGISTER,LOCK  
CMP REGISTER,#0  
JNE enter_region  
RET | return to caller; critical region entered
```

| copy lock to register and set lock to 1  
| was lock zero?  
| if it was non zero, lock was set, so loop

leave\_region:

```
MOVE LOCK,#0  
RET | return to caller
```

| store a 0 in lock

# Busy Waiting: Problems

- Waste CPU time since it sits on a tight loop
- May have unexpected effects (for priority scheduling algorithm):
  - Priority Inversion Problem
    - a high priority **waits** for a low priority to leave the critical section
    - the low priority can **never** execute since the high priority is **not blocked**.

Example:

Process H and L with **scheduling rule** that L is never scheduled while H is scheduled. Let L in CR and H is ready and wants to enter CR. Since H is ready it is given the CPU and it starts busy waiting. L will never get to run

Solution: Sleep and Wake-up

## Sleep and Wake-up

Blocking vs. Busy Waiting: Some interprocess communication (IPC) methods use blocking (not busy waiting), meaning a process stops executing instead of wasting CPU time while waiting.

Sleep and Wakeup Primitives: sleep is a system call that makes a process block (suspend execution) until another process wakes it up.

wakeup is used to resume (wake up) a specific blocked process.

# Producer Consumer Problem

- Also called bounded-buffer problem
- Two ( $m+n$ ) processes share a **common** buffer
- One ( $m$ ) of them is (are) producer(s): put(s) information in the buffer
- One ( $n$ ) of them is (are) consumer(s): take(s) information out of the buffer
- Trouble and solution
  - Producer wants to put but buffer **full**- Go to sleep and wake up when consumer takes one or more
  - Consumer wants to take but buffer **empty**- go to sleep and wake up when producer puts one or more

# Sleep and Wakeup

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                                /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        if (count == N) sleep();
        insert_item(item);
        count = count + 1;
        if (count == 1) wakeup(consumer);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        if (count == 0) sleep();
        item = remove_item();
        count = count - 1;
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);
    }
}
```

Problems  
are  
discussed  
in the  
next slide

## Producer-consumer problem

# Sleep and Wakeup: Race condition

- Unconstrained access to *count*
  - CPU is given to P just after C has **read** count to be 0 but not yet gone to sleep.
  - Result is lost wake-up signal
  - Both will sleep forever

# Lost Wakeup Signal

- Process A wants to wakeup B, who is not in sleep yet.
- Quick fix: **wakeup waiting bit**
  - When A calls `wakeup(B)` and B is not in sleep, set the bit
  - When B wants to sleep and the bit is set, B will not sleep
- Problem- When more than 2 processes? We need more bits. So the problem remains in principle.

# Semaphores(1)

- A new variable type: an integer variable to count the number of wakeups saved for future use.
- Value = 0 => No wakeups are saved
- Value > 0 => One or more wakeups are saved
- Two operations: Down and UP(similar to sleep and wakeup)
- Operation “down”: if value > 0 value-- (uses up one stored wakeup) and then continue. if value = 0, process is put to sleep without completing the down for the moment
  - Checking the value, changing it, and possibly going to sleep, is all done as an **atomic** action. So once a semaphore operation has started, no other process can access it until the operation has completed or blocked.

# Semaphores(2)

- Operation “up”:increments the value of the semaphore addressed. If one or more process were sleeping on that semaphore, one of them is chosen by the system (e.g. at random) and is allowed to complete its down
  - The operation of incrementing the semaphore and waking up one process is also **indivisible**
- No process ever blocks doing an *up*.

# Semaphores(3)

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

/\* number of slots in the buffer \*/  
/\* semaphores are a special kind of int \*/  
/\* controls access to critical region \*/  
/\* counts empty buffer slots \*/  
/\* counts full buffer slots \*/

/\* TRUE is the constant 1 \*/  
/\* generate something to put in buffer \*/  
/\* decrement empty count \*/  
/\* enter critical region \*/  
/\* put new item in buffer \*/  
/\* leave critical region \*/  
/\* increment count of full slots \*/

/\* infinite loop \*/  
/\* decrement full count \*/  
/\* enter critical region \*/  
/\* take item from buffer \*/  
/\* leave critical region \*/  
/\* increment count of empty slots \*/  
/\* do something with the item \*/

The producer-consumer problem using semaphores

# Semaphores in Producer Consumer Problem: Analysis

Initial values: empty = 3, full = 0, mutex = 1

Step	Producer Action	Consumer Action	empt y	ful l	mutex
1	P0 produces item	-	3	0	1
2	down(empty) → empty=2	-	2	0	1
3	down(mutex) → enters critical section	-	2	0	0
4	Inserts item, up(mutex)	-	2	0	1
5	up(full) → full=1	-	2	1	1
6	-	C0 calls down(full) → full=0	2	0	1
7	-	down(mutex) → enters critical section	2	0	0
8	-	Removes item, up(mutex)	2	0	1
9	-	up(empty) → empty=3	3	0	1

# Semaphores in Producer Consumer Problem: Analysis

- 3 semaphores are used
  - *full* (initially 0) for counting occupied slots
  - *Empty* (initially  $N$ ) for counting empty slots
  - *mutex* (initially 1) to make sure that P and C do not access the buffer at the same time.
- Binary Semaphores: That are initialized to one and used by 2 or more processes to ensure that only one of them enter CR at the same time
- If each process does a down just before entering CR and an up just after leaving then mutual exclusion is guaranteed
- Here 2 uses of semaphores
  - Mutual exclusion (*mutex*)
  - Synchronization (*full* and *empty*): to guarantee that certain event sequences do or do not occur

# Mutexes

- Simplified version of semaphore
- A *mutex* is a variable that can be in 1 of 2 states: locked and unlocked.
- Two procedures: mutex\_lock and mutex\_unlock
- Similar like enter\_region with a **crucial difference**

mutex\_lock:

```
TSL REGISTER,MUTEX  
CMP REGISTER,#0  
JZE ok  
CALL thread_yield  
JMP mutex_lock
```

ok: RET | return to caller; critical region entered

| copy mutex to register and set mutex to 1  
| was mutex zero?  
| if it was zero, mutex was unlocked, so return  
| mutex is busy; schedule another thread  
| try again later

mutex\_unlock:

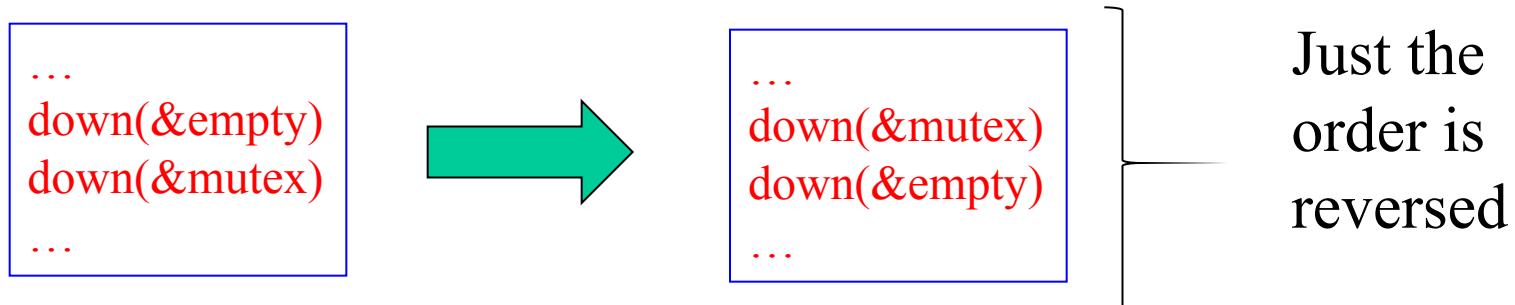
```
MOVE MUXTEX,#0  
RET | return to caller
```

| store a 0 in mutex

Implementation of *mutex\_lock* and *mutex\_unlock*

# Semaphores: “Be Careful”

- Suppose the following is done in P’s code



- If buffer **full** P would block due to `down(&empty)` with `mutex = 0`.
- So now if C tries to access the buffer, it would block too due to its `down(&mutex)`.
- Both processes would stay blocked **forever**:  
**DEADLOCK**

# Revisiting semaphores!

- Semaphores are very “low-level” primitives
  - Users could easily make small errors
  - Very **difficult to debug**
- Also, we seem to be using them in two ways
  - For mutual exclusion, the “real” abstraction is a **critical section**
  - But for Synchronization threads “**communicate**” using semaphores
- Simplification: Provide concurrency support in compiler
  - Monitors

\*

# Monitors (1)

- A higher level synchronization primitive
- A **collection** of procedures, variables and data structures grouped in a special kind of **module** or **package**.

**monitor** *example*

**integer** *i*;  
**condition** *c*;

**procedure** *producer*( );

.

.

.

**end**;

**procedure** *consumer*( );

.

.

.

**end**;

**end monitor**;

Example of a monitor

# Monitors(2)

- Only one process can be active in a monitor at any instant
- Monitors are programming language construct, so the compiler knows they are special and can handle calls to monitor procedures differently from other calls.
- Because the compiler, not the programmer, is arranging the mutual exclusion it is safer
- We also need a way to block and wakeup: Wait and Signal (done on a condition variables)

# Monitors(3)

- *wait* is called on some condition variables:
  - Calling process is **blocked** and another process that had been previously prohibited from entering the monitor is **allowed** to enter now.
- *signal* is called on some condition variable:
  - A process waiting on *that CV* is given the chance to get up.
  - Who should run? Caller or awakened one?

**Alternative#1:** Let newly awakened process to run suspending the caller.

**Alternative#2:** Process doing a signal must exit the monitor immediately i.e. signal statement may appear only as the final statement in a monitor procedure.

**Alternative#3:** Let the caller run and when it exits the monitor then the waiting process is allowed to start.

**Note: If more than one processes are waiting on full, one of them is scheduled.**

# Monitors (4)

```
monitor ProducerConsumer
    condition full, empty;
    integer count;
    procedure insert(item: integer);
begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
end;
function remove: integer;
begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
end;
count := 0;
end monitor;

procedure producer;
begin
    while true do
begin
    item = produce_item;
    ProducerConsumer.insert(item)
end
end;
procedure consumer;
begin
    while true do
begin
    item = ProducerConsumer.remove;
    consume_item(item)
end
end;
```

- Outline of producer-consumer problem with monitors
  - only one monitor procedure active at one time
  - buffer has  $N$  slots

- **Problems with monitors and semaphores**
  - Semaphores are too low level
  - Monitors are not usable except in a few programming languages
  - Designed to work in an environment having access to a **common** memory
  - Doesn't allow information exchange among machines
  - None of them would work in a distributed systems (why?) consisted of multiple CPUs, each with its **own** private memory.

# Message Passing

- solution to the problem of semaphores and monitors w.r.t distributed systems
- A method of IPC that uses two primitives
  - send and receive: system calls.
  - send(destination, &message)
  - receive(source, &message)
  - If no message is available:
    - The receiver can **block** until one arrives
    - Return immediately with an error code

# Message Passing

- Challenges:
  - Messages can be lost by the network.
  - Acknowledgement and retransmission issue.
  - Process naming.
  - Authentication.
  - Performance issue.

# Producer Consumer with Message Passing

## Assumptions:

- All messages are of same size
- Messages sent but not yet received are automatically buffered

```
#define N 100                                     /* number of slots in the buffer */  
  
void producer(void)  
{  
    int item;  
    message m;  
  
    while (TRUE) {  
        item = produce_item();  
        receive(consumer, &m);  
        build_message(&m, item);  
        send(consumer, &m);  
    }  
}  
  
void consumer(void)  
{  
    int item, i;  
    message m;  
  
    for (i = 0; i < N; i++) send(producer, &m); /* send N empties */  
    while (TRUE) {  
        receive(producer, &m);  
        item = extract_item(&m);  
        send(producer, &m);  
        consume_item(item);  
    }  
}
```

```
/* message buffer */  
  
/* generate something to put in buffer */  
/* wait for an empty to arrive */  
/* construct a message to send */  
/* send item to consumer */
```

The producer-consumer problem with N messages

# Many Variants of Message Passing

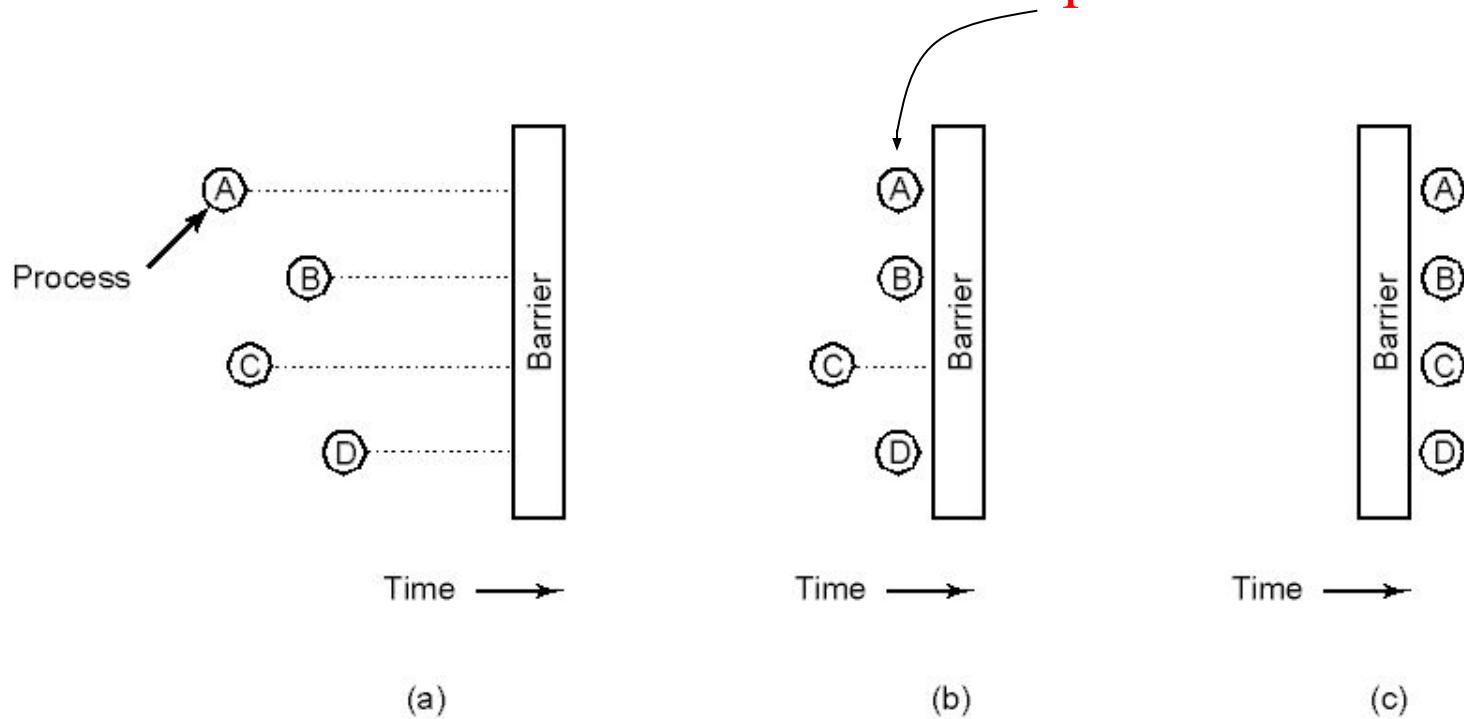
- How messages are addressed?
  - Assign each process a unique address and have messages be addressed to processes
  - Invent and use a new data structure, **mailbox**-a place to buffer a certain number of messages. So here mailboxes have addresses
- Buffering?
  - Mailbox is a kind of buffering
  - Other extreme is to eliminate buffering- if *send* is done first then sender is **blocked** until receive happens and vice versa (Easier to implement than buffering)

# Barriers

- A synchronizing mechanism intended for groups of processes
- Some applications are divided into phases and have the rule that no process may proceed into the next phase until all processes are ready to proceed to the next phase.
- This can be done by placing a barrier at the end of each phase.
- When a process reaches the barrier, it is blocked until all processes have reached the barrier.

# Barriers

Executes a barrier primitive



- Use of a barrier
  - processes approaching a barrier
  - all processes but one blocked at barrier
  - last process arrives, all are let through

# Barriers

```
// Pseudocode Example: Shared variable  
barrier(count = 3);
```

```
Thread_1() {  
    compute_part1();  
    wait_at_barrier();  
    combine_results();  
}
```

```
Thread_2() {  
    compute_part2();  
    wait_at_barrier();  
    combine_results();  
}
```

```
Thread_3() {  
    compute_part3();  
    wait_at_barrier();  
    combine_results();  
}
```

# Semaphore Usage

- A lock permits only **1** process/thread to be in Critical Section.
- What if we want maximum **m** process/thread to be in the critical section simultaneously ?
- Semaphore from railway analogy
  - Here is a semaphore **initialized to 2** for resource control:

