

Chapter 3

Memory Management


Main Memory (RAM)

- Main memory (RAM) is an important resource that must be carefully managed.
- A program resides on a disk as binary executable file
- To be executed the program must be brought into RAM
- The CPU fetches instructions from RAM according to the value of the PC
- The instruction operands may be needed to be fetched from RAM
- After execution the results may be stored back to RAM

Memory Management

- Ideally programmers want memory that is
 - private
 - large
 - fast
 - non volatile
 - Cheap

Computer Hardware Review



- It is the job of the operating system to
 - **abstract** this hierarchy into a useful model
 - and then **manage** the abstraction.

Memory Management

- The part of the operating system that manages the memory hierarchy is called the *memory manager*.
 - Keep track of used memory
 - Allocate memory to processes
 - Deallocate it when they are done


Objective

- In this chapter we will study
 - how operating systems create **abstractions** from **memory**
 - and how they **manage** them.
- The focus will be on the **programmer's model** of main memory
- Memory Abstraction: the view/illusion of the memory presented to the programmer by the OS


No Memory Abstraction

- The simplest memory abstraction
- Every program simply saw the physical memory
 - MOV REG1, 1000
 - move the contents of physical memory location 1000 to REGISTER1
- Model of memory presented to the programmer is simply physical memory
- Not possible to have 2 running programs in memory at the same time


No Memory Abstraction



(a)



(b)



(c)

Three simple ways of organizing memory with an operating system and one user process



Running Multiple Programs Without a Memory Abstraction

- Used in the early models of the IBM 360.
- Memory can be divided into 2-KB blocks and each one was assigned a 4-bit protection key held in special registers inside the CPU.
- The PSW (Program Status Word) also contained a 4-bit key.
- The hardware trapped any attempt by a running process to access memory with a protection code different from the PSW key.


Running Multiple Programs Without a Memory Abstraction

- Relocation problem
- Use of static relocation


Exposing physical memory to processes is not a good idea



(a)



(b)



(c)

Conclusion

- exposing physical memory to processes has several major drawbacks.
 - if user programs can address every byte of memory, they can easily trash the operating system intentionally or by accident
 - it is difficult to have multiple programs running at once

Conclusion

- Two problems have to be solved to allow multiple applications to be in memory at the same time without interfering with each other
 - protection
 - Relocation
- SOLUTION
 - invent a new **abstraction** for memory: the address space

A Memory Abstraction: ADDRESS SPACES


- An address space is the range of addresses that a **process** can use to address memory.
 - Just as the process concept creates a kind of *abstract CPU* to run programs,
 - the address space creates a kind of *abstract memory* for programs to live in.
- Each process has its own **independent** address space

A Memory Abstraction: ADDRESS SPACES


- how to give each program its own address space??
 - address 28 in one program means a different physical location than address 28 in another program

Base and Limit Registers

- Base and limit registers can be used to give each process a separate address space
- CPU hardware automatically adds the base value to the address generated by the process
 $JMP\ 28 \Rightarrow JMP\ 16412$
- Dynamic relocation.



Base and Limit Registers



Disadvantage


- an addition and a comparison on every memory reference.

Memory Overload


- In practice, the total amount of RAM needed by all the processes is often much more than can fit in memory.
- keeping all processes in memory all the time requires a huge amount of memory and cannot be done if there is insufficient memory.
- Two approaches to solve
 - Swapping
 - Virtual Memory

Swapping (1)

- Bringing in each process in its **entirety**,
- running it for a while
- then putting it back on the disk



Swapping




Memory allocation changes as

- processes come into memory
- leave memory

Shaded regions are unused memory

Contiguous Allocation in Swapping


- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)



Swapping


- When swapping creates multiple holes in memory, it is possible to **combine** them all into one **big** one by moving **all the processes** downward as far as possible.
- This technique is known as **memory compaction**

Swapping



- Allocating space for growing data segment
- Allocating space for growing stack & data segment


Memory Layout for Process



Managing Free Memory

- Operating system maintains information about
 - Allocated memory
 - Free memory (holes)
- there are two ways to keep track of memory usage:
 - bitmaps
 - free lists

Memory Management with Bit Maps




- memory is divided into allocation units
- Corresponding to each allocation unit is a bit in the bitmap

Bitmaps


- The good-compact way to keep track of memory
- The bad-need to search memory for k consecutive zeros to bring in a file k units long
- Units can be bits or bytes or.....

Memory Management with Linked Lists

- maintain a linked list of allocated and free memory segments,
 - where a segment either contains a process or is an empty hole between two processes



Memory Management with Linked Lists



Four neighbor combinations for the terminating process X
Here the segment list is kept sorted by **address**

- Might want to use doubly linked lists for to merge holes more easily

Linked Lists

- Need algorithms to fill in the holes in memory
 - First fit
 - Best fit
 - Next fit
 - Worst fit
 - Quick fit

The fits

First-Fit Algorithm

- **Idea:**
Allocate the **first hole** that is **large enough** to accommodate the process.
- The OS searches the memory from the **beginning** and stops as soon as it finds a suitable hole.
- **Advantages:**
 - Simple and fast (minimal search time).
 - Low overhead.
- **Disadvantages:**
 - May leave **many small, unusable holes** near the beginning of memory (**external fragmentation**).

The fits

Best-Fit Algorithm

- **Idea:**

Allocate the **smallest hole that is big enough** for the process.

- **Advantages:**

- Reduces wasted space in large holes.

- **Disadvantages:**

- **Slow**, since it must check **all holes** to find the smallest that fits.

The fits

Next-Fit Algorithm

- **Idea:**

Similar to **First-Fit**, but instead of always starting from the beginning of memory, it starts searching from **the point of the last allocation** and wraps around circularly.
- **Advantages:**
 - More evenly distributes memory use.
 - Prevents the beginning of memory from getting cluttered.
- **Disadvantages:**
 - Slightly more overhead than first-fit.

The fits

Worst-Fit Algorithm

- **Idea:**
Allocate the **largest available hole**, leaving the biggest remaining space for future allocations.
- **Advantages:**
 - May reduce the number of small holes.
- **Disadvantages:**
 - Often **wastes large memory blocks**.

The fits


Quick-Fit Algorithm

- **Idea:**

Maintains **separate lists** of holes for **commonly requested sizes** (e.g., 4 KB, 8 KB, 16 KB). When a process arrives, the OS directly looks into the corresponding list.
- **Advantages:**
 - **Very fast** allocation for frequent sizes.
- **Disadvantages:**
 - Requires **extra bookkeeping** to maintain multiple lists.

Memory Management with Linked Lists

- Memory allocation algorithms : First fit, next fit, best fit, worst fit, quick fit



The fits

- Conclusion: the fits couldn't out-smart the un-knowable distribution of hole sizes
- The extra work to deal with something which you can't predict failed to produce good results
- Shocking!

Swapping: Conclusion

- Swapping requires that an **entire** process be in a **single contiguous** section of memory before it can execute.
- Swapping is not an attractive option as transferring data with disk is a slow operation