


Chapter 3


Memory Management

Virtual Memory

- a technique that allows the execution of processes that are **not completely in memory.**
- **abstracts** main memory into an *extremely large, uniform* array of storage as viewed by the programmers
- major advantages
 - programs can be larger than physical memory.
 - works just fine in a multiprogramming system, with pieces of **many** programs in memory at once.
 - While a program is waiting for **piece of itself to be read in**, the **CPU** can be given to another process.

Paging : A Technique used in Virtual Memory System






CPU: Central Processing Unit

MMU: Memory Management Unit

TLB: Translation lookaside buffer

Paging : A Technique used in Virtual Memory System

- program-generated addresses are called virtual addresses and form the virtual address space
- The virtual address space is divided into fixed-size units called pages.
 - Each page is a contiguous range of addresses
 - These pages are mapped onto physical memory
- The corresponding units in the physical memory are called page frames.
- Transfers between RAM and disk are always in whole pages



Address translation


virtual address $\xrightarrow{\text{page table}}$ *physical address*

Page 0 $\xrightarrow{\text{map to}}$ *Frame 2*

0_{virtual} $\xrightarrow{\text{map to}}$ 8192_{physical}

20500_{vir} $\xrightarrow{\text{map to}}$ 12308_{phy}
 $(20k + 20)_{\text{vir}}$ $\xrightarrow{\text{map to}}$ $(12k + 20)_{\text{phy}}$

Page Fault



What happens if the program references an unmapped addresses???

MOV REG, 32780 ?


Page fault & swapping

Page Fault


- If the MMU notices that the requested **virtual** page is unmapped it causes the CPU to trap to the **operating system**.
- This trap is called a **page fault**.
- The **operating system**
 - picks a little-used page frame and writes its contents back to the disk (if needed).
 - fetches the page just referenced into the page frame just freed
 - changes the map, and restarts the trapped instruction.

Paging Request

Request Page 3





Virtual Memory Stored **on Disk**



Page Table
VM Frame


3	1
	2
	3
	4

Real Memory




Paging

Request Page 1




Real Memory




Paging

Request Page 6




Real Memory



Paging

Request Page 2




Paging

Request Page 8

Swap page 1 out


Memory



Page Table
VM Frame

VM	Frame
3	1
	2
6	3
2	4


Real Memory




Paging

Request Page 8

Swap page 1 out; load page
8



Real Memory



Inside the MMU: Page Table

Address generated by CPU is divided into:


Page number(p): an index into a *page table*

Page offset(d): to be copied into memory


Given *logical address space* 2^m and *page size* 2^n ,


$$\text{number of pages} = \frac{2^m}{2^n} = 2^{m-n}$$

Example: addressing to 001000000000100




$$\text{page number} = 0010 = 2, \quad \text{page offset} = 0000000000100$$





Page Tables Entry

- Entry for each virtual page in the page table



Paging

- 2 major issues
 - The mapping from virtual address to physical address must be fast.
 - If the virtual address space is large, the page table will be large.